

Tree-based Methods¹

Shaobo Li

University of Kansas

¹Partially based on Hastie, et al. (2009) ESL, and James, et al. (2013) ISLR

Overview

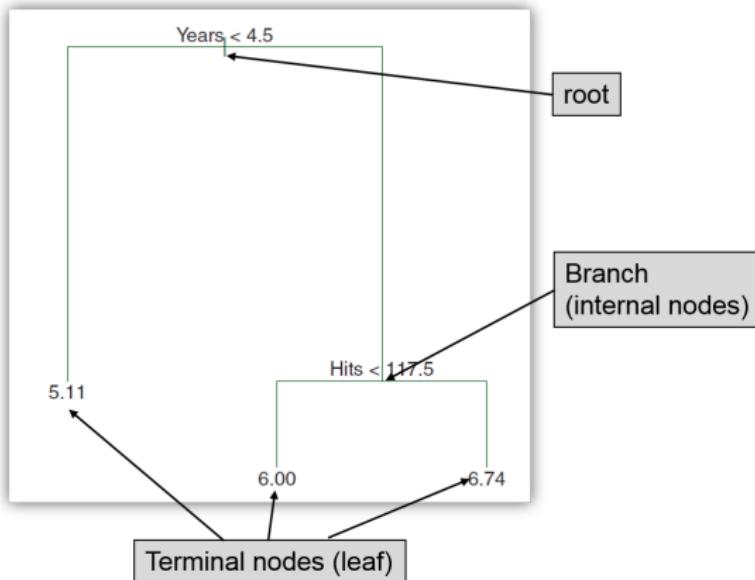
- Supervised learning
- Simple for interpretation
- Nonparametric method

We will talk about

- Classification and Regression Tree (CART)
- Bagging, Random Forest, and Boosting

A Simple Example of Regression Tree

- Predicting baseball player's log-salary using
 - X1: number of years he has played in the major leagues;
 - X2: number of hits he made in previous year



Prediction

- The values in terminal nodes
- Top-down
- At each splitting point, go left if TRUE

Data Partition — Stratification

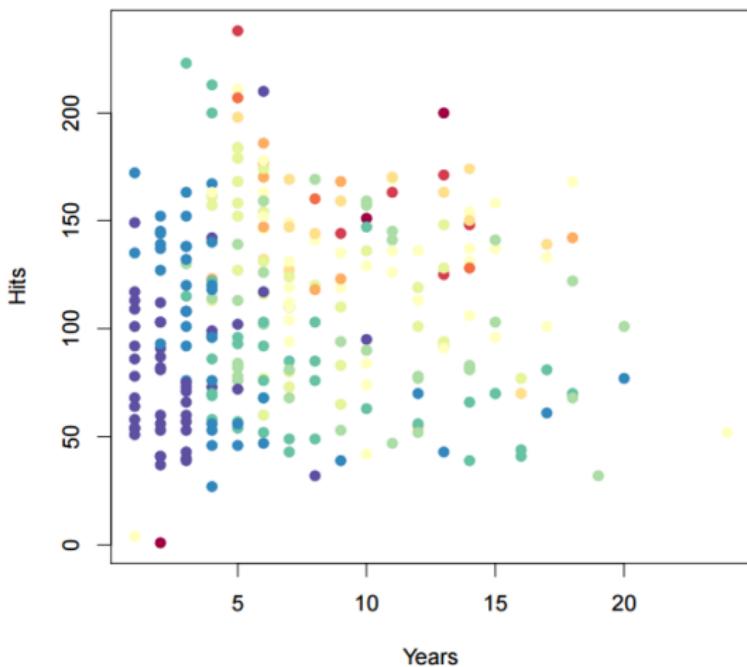
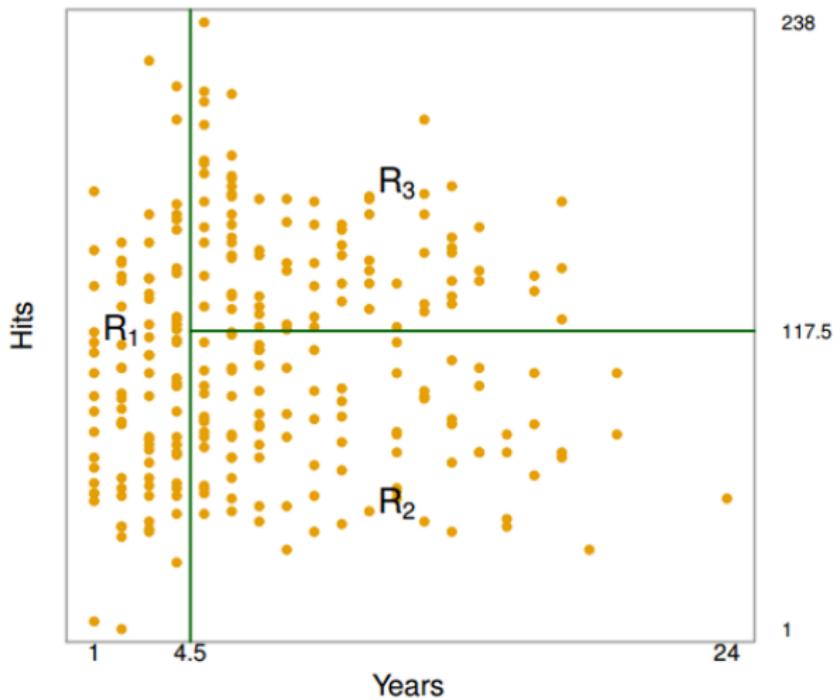


Figure: Salary is color-coded from low (blue, green) to high (yellow, red)

Data Partition — Stratification



A few technical question

- How to determine the optimal splitting point?
- What criteria should we use?
- When to stop growing the tree?

Optimal split (for a regression tree)

- For each split, we try to determine two regions R_1 and R_2 by finding the optimal splitting point $X_j = s$, that is

$$R_1 = \{\text{Data} | X_j \leq s\} \quad \text{and} \quad R_2 = \{\text{Data} | X_j > s\}$$

This is done by scanning all possible splitting point for each X .

- Such optimal splitting point minimizes residual sum squares

$$\sum_{i \in R_1} (y_i - \bar{y}_{R_1})^2 + \sum_{i \in R_2} (y_i - \bar{y}_{R_2})^2$$

where \bar{y}_{R_1} and \bar{y}_{R_2} are average of Y in each region.

- Repeat this process for each region until the decrease of RSS is not significant.

Stopping Strategy — Cost-Complexity Pruning

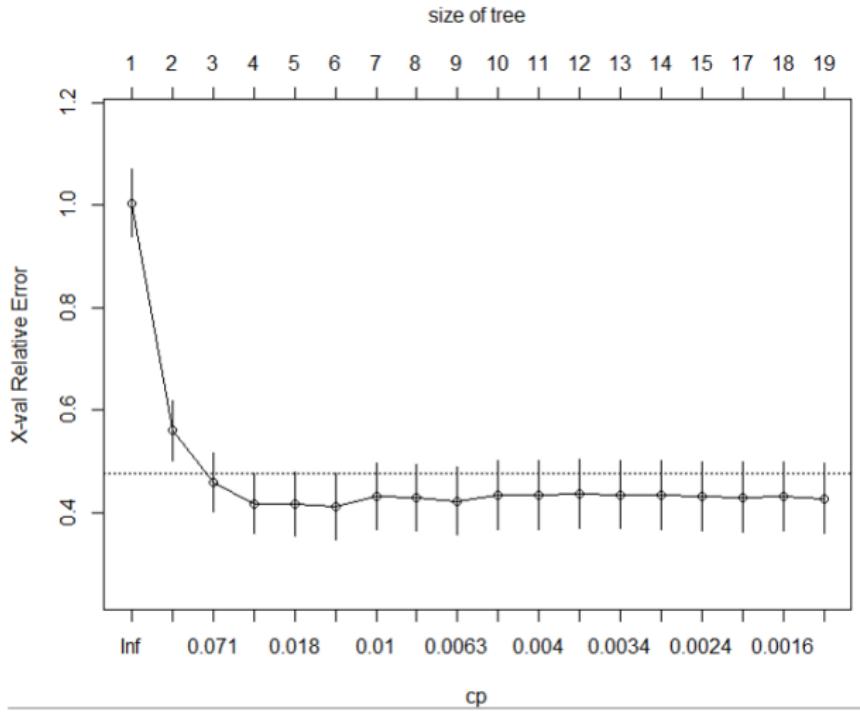
- To avoid overfitting!
- Suppose T_0 is a very large tree (overfitted tree). Denote T as a subtree of T_0 , and $|T|$ as the size of that tree (number of terminal nodes). We minimize the cost complexity criterion

$$C_\alpha(T) = \sum_{m=1}^{|T|} RSS_m + \alpha |T|,$$

where RSS_m is residual sum squares in region (terminal) m .

- Larger α penalizes the tree size ($|T|$) more
 - Recall how do we use AIC/BIC for variable selection
 - Every α corresponds to a unique optimal tree T_α .
- We use **cross-validation** to choose the optimal α .

Pruning a Tree in R

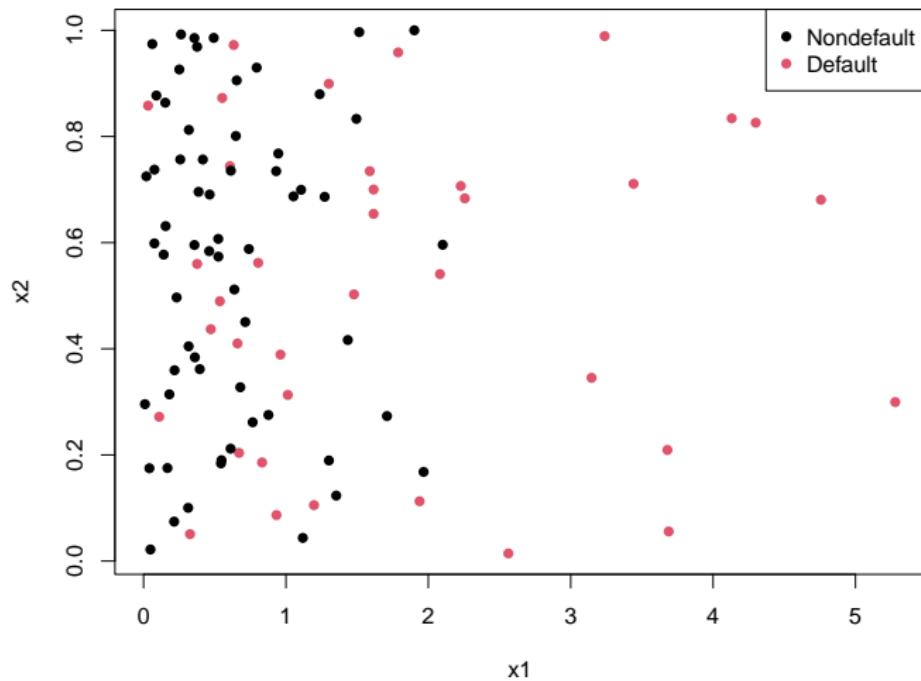


- cp is the complexity parameter α we have discussed.

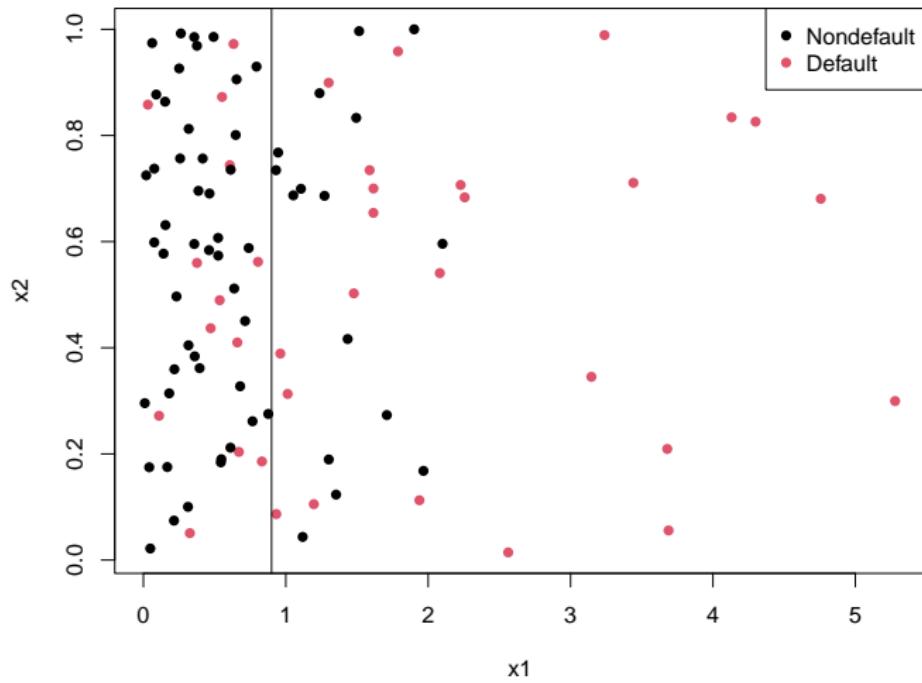
Classification Tree

- Outcome is categorical variable
- The same idea as for regression tree
- The **only difference is the splitting criteria**
- What should we use for the criteria?

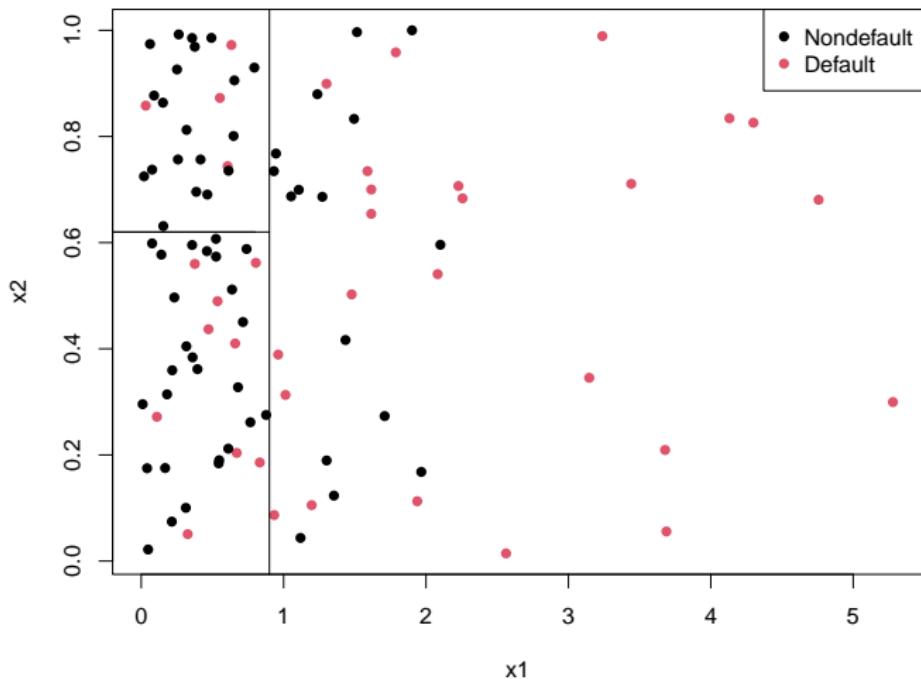
Illustration



Illustration



Illustration



Splitting Criteria for Classification Tree

- Misclassification rate
- Gini index

$$\text{Gini index} = \sum_{k \neq k'} \hat{p}_k \hat{p}_{k'} = \sum_{k=1}^K \hat{p}_k (1 - \hat{p}_k)$$

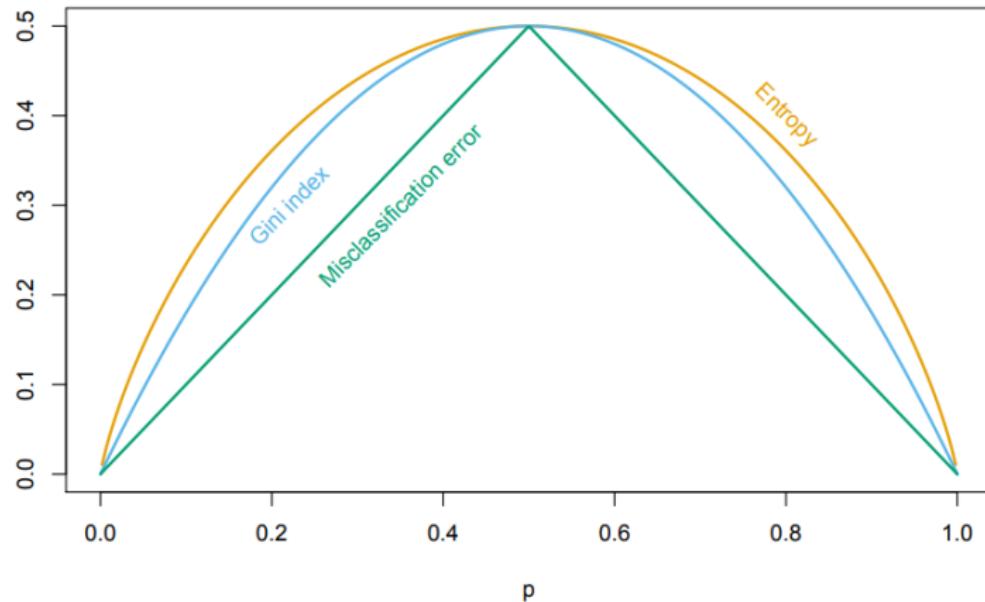
where \hat{p}_k is the proportion of k category observation in one node.

- Entropy

$$\text{Entropy} = - \sum_{k=1}^K \hat{p}_k \log \hat{p}_k$$

For each split, gini index (or entropy) is calculated for two regions, and the weighted average is used as the criterion to be *minimized*.

Illustration



Example

- Suppose a node contains binary outcome (0 or 1) with each 400 observations.
- Split strategy 1: (300, 100) and (100, 300)
- Split strategy 2: (400, 200) and (0, 200)
- What is the misclassification rate, Gini index, and Entropy for each split strategy?

Categorical Predictor

- If a categorical predictor has q possible values, how many possible splits?
- Computationally infeasible for large q .
- Solution for binary outcome:
 - Sort the q classes according to the proportion of "1" in each of the q classes.
 - Split the ordered variable as if it is numeric variable.
 - This split is the optimal in terms of Gini index and Entropy.
- But we should avoid categorical variable which has too many categories. Why?

Asymmetric Cost

- Similar idea as we discussed in last lecture
- Define a loss matrix with different weights in misclassified cells.

	Pred=1	Pred=0
True=1	0	w_0
True=0	w_1	0

- Loan application example (1=default): $w_1 < w_0$
- The weights are incorporated during each split

Bagging

- Limitation of a single tree: high variance, unstable
 - Change of points in the sample would lead to different split.
 - Current split heavily relies on previous splits.
- Bagging — Bootstrap aggregation

$$\hat{f}_{\text{bag}}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}_b(x)$$

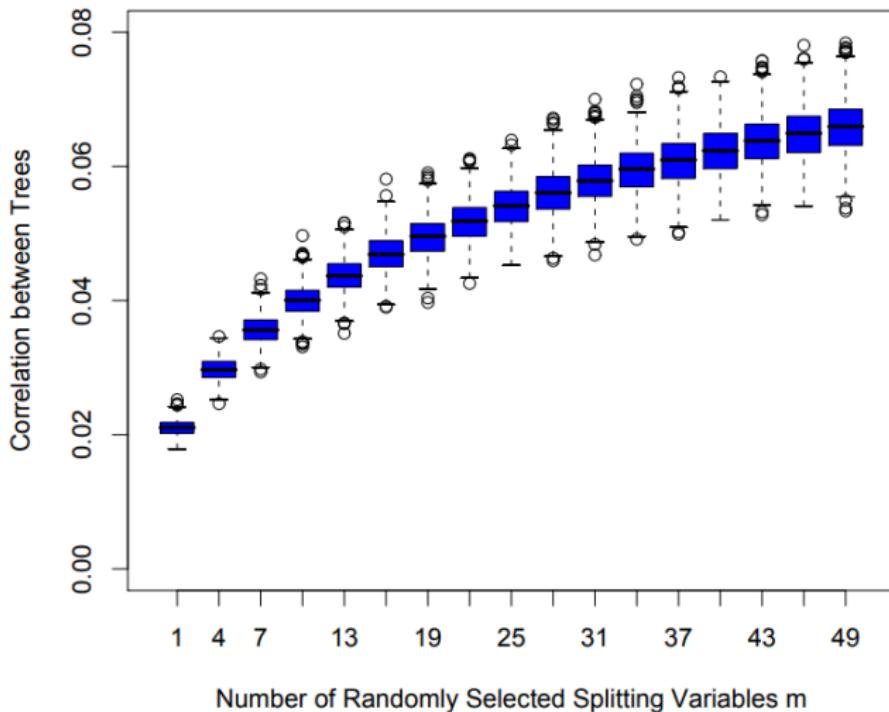
- Leo Breiman (1996)
- Fit many trees on bootstrap samples, and then take average.
- Variance can be significantly reduced
- For classification tree, how to take average?

Out-of-bag (OOB) Prediction

- Similar to cross-validation
- A type of testing error — to prevent overfitting
- How OOB is computed:
For each observation i :
 - 1 Identify all bootstrap samples that did not include i .
 - 2 Predict i 's observation using all the trees based on those bootstrap samples.
 - 3 Average (for regression) or majority vote (for classification) across those trees to get the OOB prediction for observation i .

- Improvement of Bagging
 - The trees in Bagging are correlated to each other
 - It may add variance and prediction error
- Goal: decorrelate a series of trees while maintaining the strength of ensemble.
- How RF works: randomly choose m predictors as candidate splitting variables for each split
- See the inventor's (Leo Breiman) [website](#) for more details.

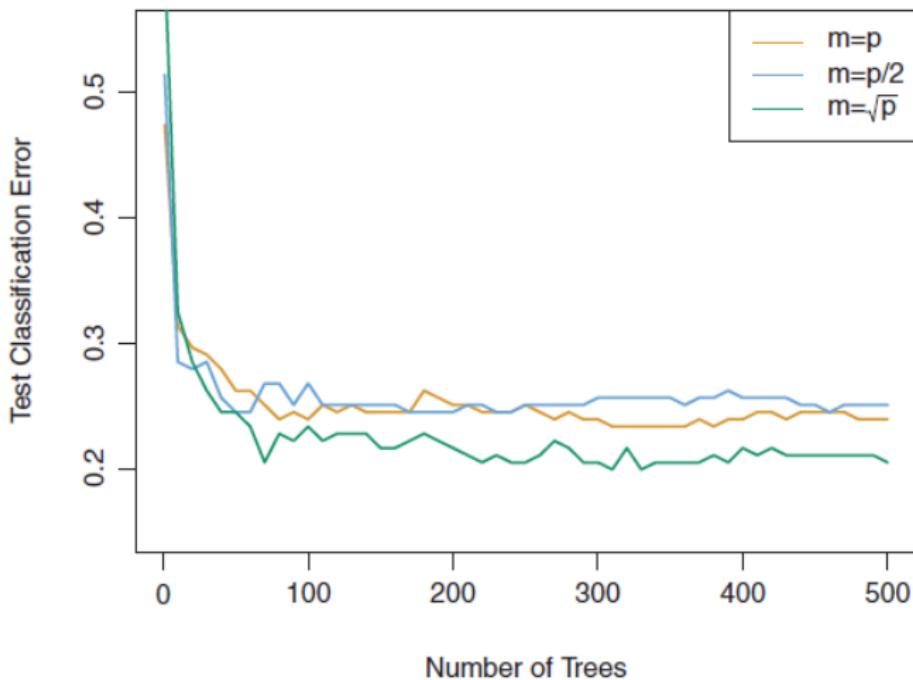
Illustration of De-correlation



Choice of m

- Smaller m decreases the correlation, but also decreases strength
- Larger m increases the correlation, but also increases strength
- Use OOB error to determine the optimal m
- By default, $m \approx \sqrt{p}$ for classification and $p/3$ for regression

Change of Testing Error across m



Gradient Boosting

- Ensemble method
- Based on small trees (weak learner)
- Additive – adding up all small trees
- An illustration
- Pros:
 - Powerful algorithm, high prediction accuracy
 - Small variance (advantage of ensemble methods)
- Cons:
 - Computational expensive for large data
 - Lack of interpretability
 - May overfit data

Boosting Regression Tree

- Additive: sequentially grow the tree
- Combination of many small trees
- Boosting for regression tree
 - 1 Set $\hat{f}(x) = 0$, and $r_i = y_i$
 - 2 For $b = 1, \dots, B$, repeat:
 - Fit a small tree $\hat{f}^{(b)}(x)$ with d splits based on the training sample
 - Update residual $r_i \leftarrow r_i - \lambda \hat{f}^{(b)}(x)$, and treat this residual as new response, where λ is shrinkage parameter
 - Update the tree $\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^{(b)}(x)$
 - 3 Output the boosted regression tree

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^{(b)}(x)$$

Gradient Boosting Machine (GBM)

- A more generalized version.
- User supplied differentiable loss function (e.g. squared loss)

$$\sum_{i=1}^n L(y_i, f(x_i)) = \sum_{i=1}^n (y_i - f(x_i))^2$$

- The negative **gradient** for each $L(y_i, f(x_i))$ is essentially residual or pseudo residual.
- Then fit a new regression tree to approximate the gradient (fit the residual).
- For binary classification, the loss can be binomial deviance (equivalent to cross entropy, see ESL p346), but the weak learner in each iteration is still a regression tree that approximate the gradient of the deviance ($y_i - p_i$).

The algorithm (ESL p387)

Algorithm 10.4 Gradient Boosting for K -class Classification.

1. Initialize $f_{k0}(x) = 0$, $k = 1, 2, \dots, K$.

2. For $m=1$ to M :

(a) Set

$$p_k(x) = \frac{e^{f_k(x)}}{\sum_{\ell=1}^K e^{f_\ell(x)}}, \quad k = 1, 2, \dots, K.$$

(b) For $k = 1$ to K :

i. Compute $r_{ikm} = y_{ik} - p_k(x_i)$, $i = 1, 2, \dots, N$.

ii. Fit a regression tree to the targets r_{ikm} , $i = 1, 2, \dots, N$, giving terminal regions R_{jkm} , $j = 1, 2, \dots, J_m$.

iii. Compute

$$\gamma_{jkm} = \frac{K-1}{K} \frac{\sum_{x_i \in R_{jkm}} r_{ikm}}{\sum_{x_i \in R_{jkm}} |r_{ikm}|(1 - |r_{ikm}|)}, \quad j = 1, 2, \dots, J_m.$$

iv. Update $f_{km}(x) = f_{k,m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jkm} I(x \in R_{jkm})$.

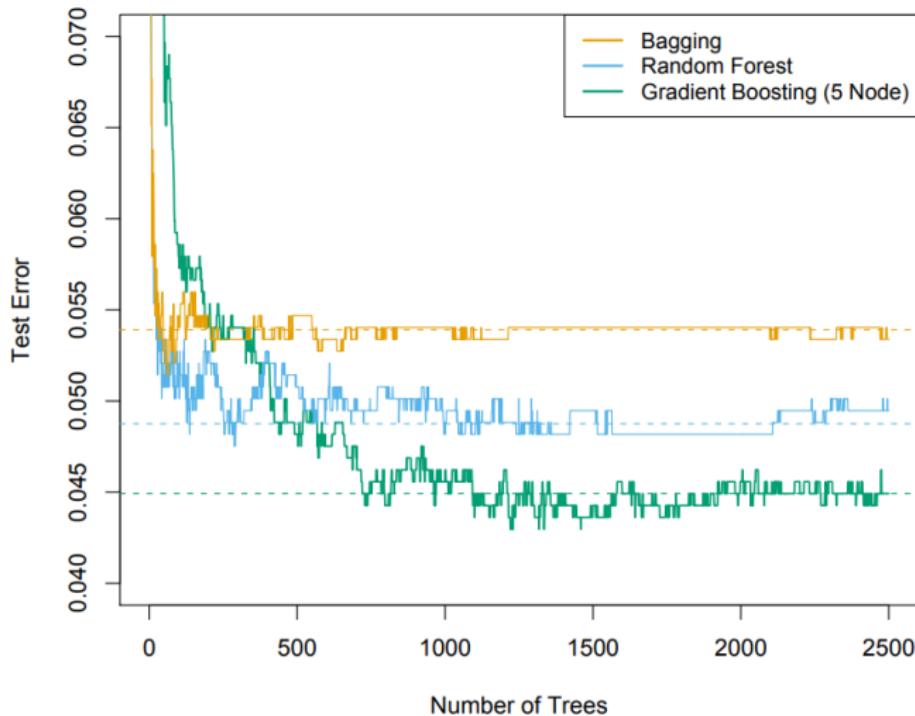
3. Output $\hat{f}_k(x) = f_{kM}(x)$, $k = 1, 2, \dots, K$.

AdaBoost — Boosting for Classification

- Originally designed for classification $Y = \{-1, 1\}$
- Algorithm:
 - 1 Assign equal weights to all observations $w_i = 1/N$
 - 2 For $b = 1, \dots, B$, repeat:
 - Fit a classifier $G_b(x)$ to the training sample using weights w_i
 - Compute weighted misclassification error
$$\text{err}_b = \frac{\sum_{i=1}^N w_i \mathbb{I}(y_i \neq G_b(x_i))}{\sum_{i=1}^N w_i}$$
 - Compute $\alpha_b = \log((1 - \text{err}_b)/\text{err}_b)$
 - Update weights $w_i \leftarrow w_i \exp\{\alpha_b \mathbb{I}(y_i \neq G_b(x_i))\}$, and normalize such that the sum equals to 1.
 - 3 Output $G(x) = \text{sign}\{\sum_{b=1}^B \alpha_b G_b(x)\}$.

Here is the [StatQuest](#) video that may help you to understand the idea.

Performance Comparison



XGBoost – Extreme Gradient Boosting

- Scalable tree boosting
- Recently developed by Tianqi Chen in 2016.
- One of the best supervised learning algorithms ([Machine Learning Challenge Winning Solutions](#))

General idea

We omit the mathematical details here. For more detail, please see the [official tutorial](#), and [original paper](#).

- Similar to traditional gradient tree boosting.
- Added regularization which prevents overfitting.
- Optimized the utility of computing power.